


Efficient Transformers

Angela Fan

Standard Neural Networks


TRAINING TIME



- Overparameterized

Standard Neural Networks


TRAINING TIME



- Overparameterized
- Redundant

Standard Neural Networks


TRAINING TIME



- Overparameterized
- Redundant
- Overfitting

Standard Neural Networks

TRAINING TIME



- Overparameterized
- Redundant
- Overfitting
- Too Large for Practical Applications

Techniques for Smaller Networks

- Train Smaller Network from Scratch

Techniques for Smaller Networks

- Train Smaller Network from Scratch
- Sparsity Inducing Training

Techniques for Smaller Networks

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation

Techniques for Smaller Networks

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing

Techniques for Smaller Networks

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization

Techniques for Smaller Networks

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization
- More efficient architectures

Caveat: this is a **brief** overview focused on **Transformers**

What to think about when talking about efficiency?

- Training Time
- Inference Time
- Model size
- Energy

What to think about when talking about efficiency?

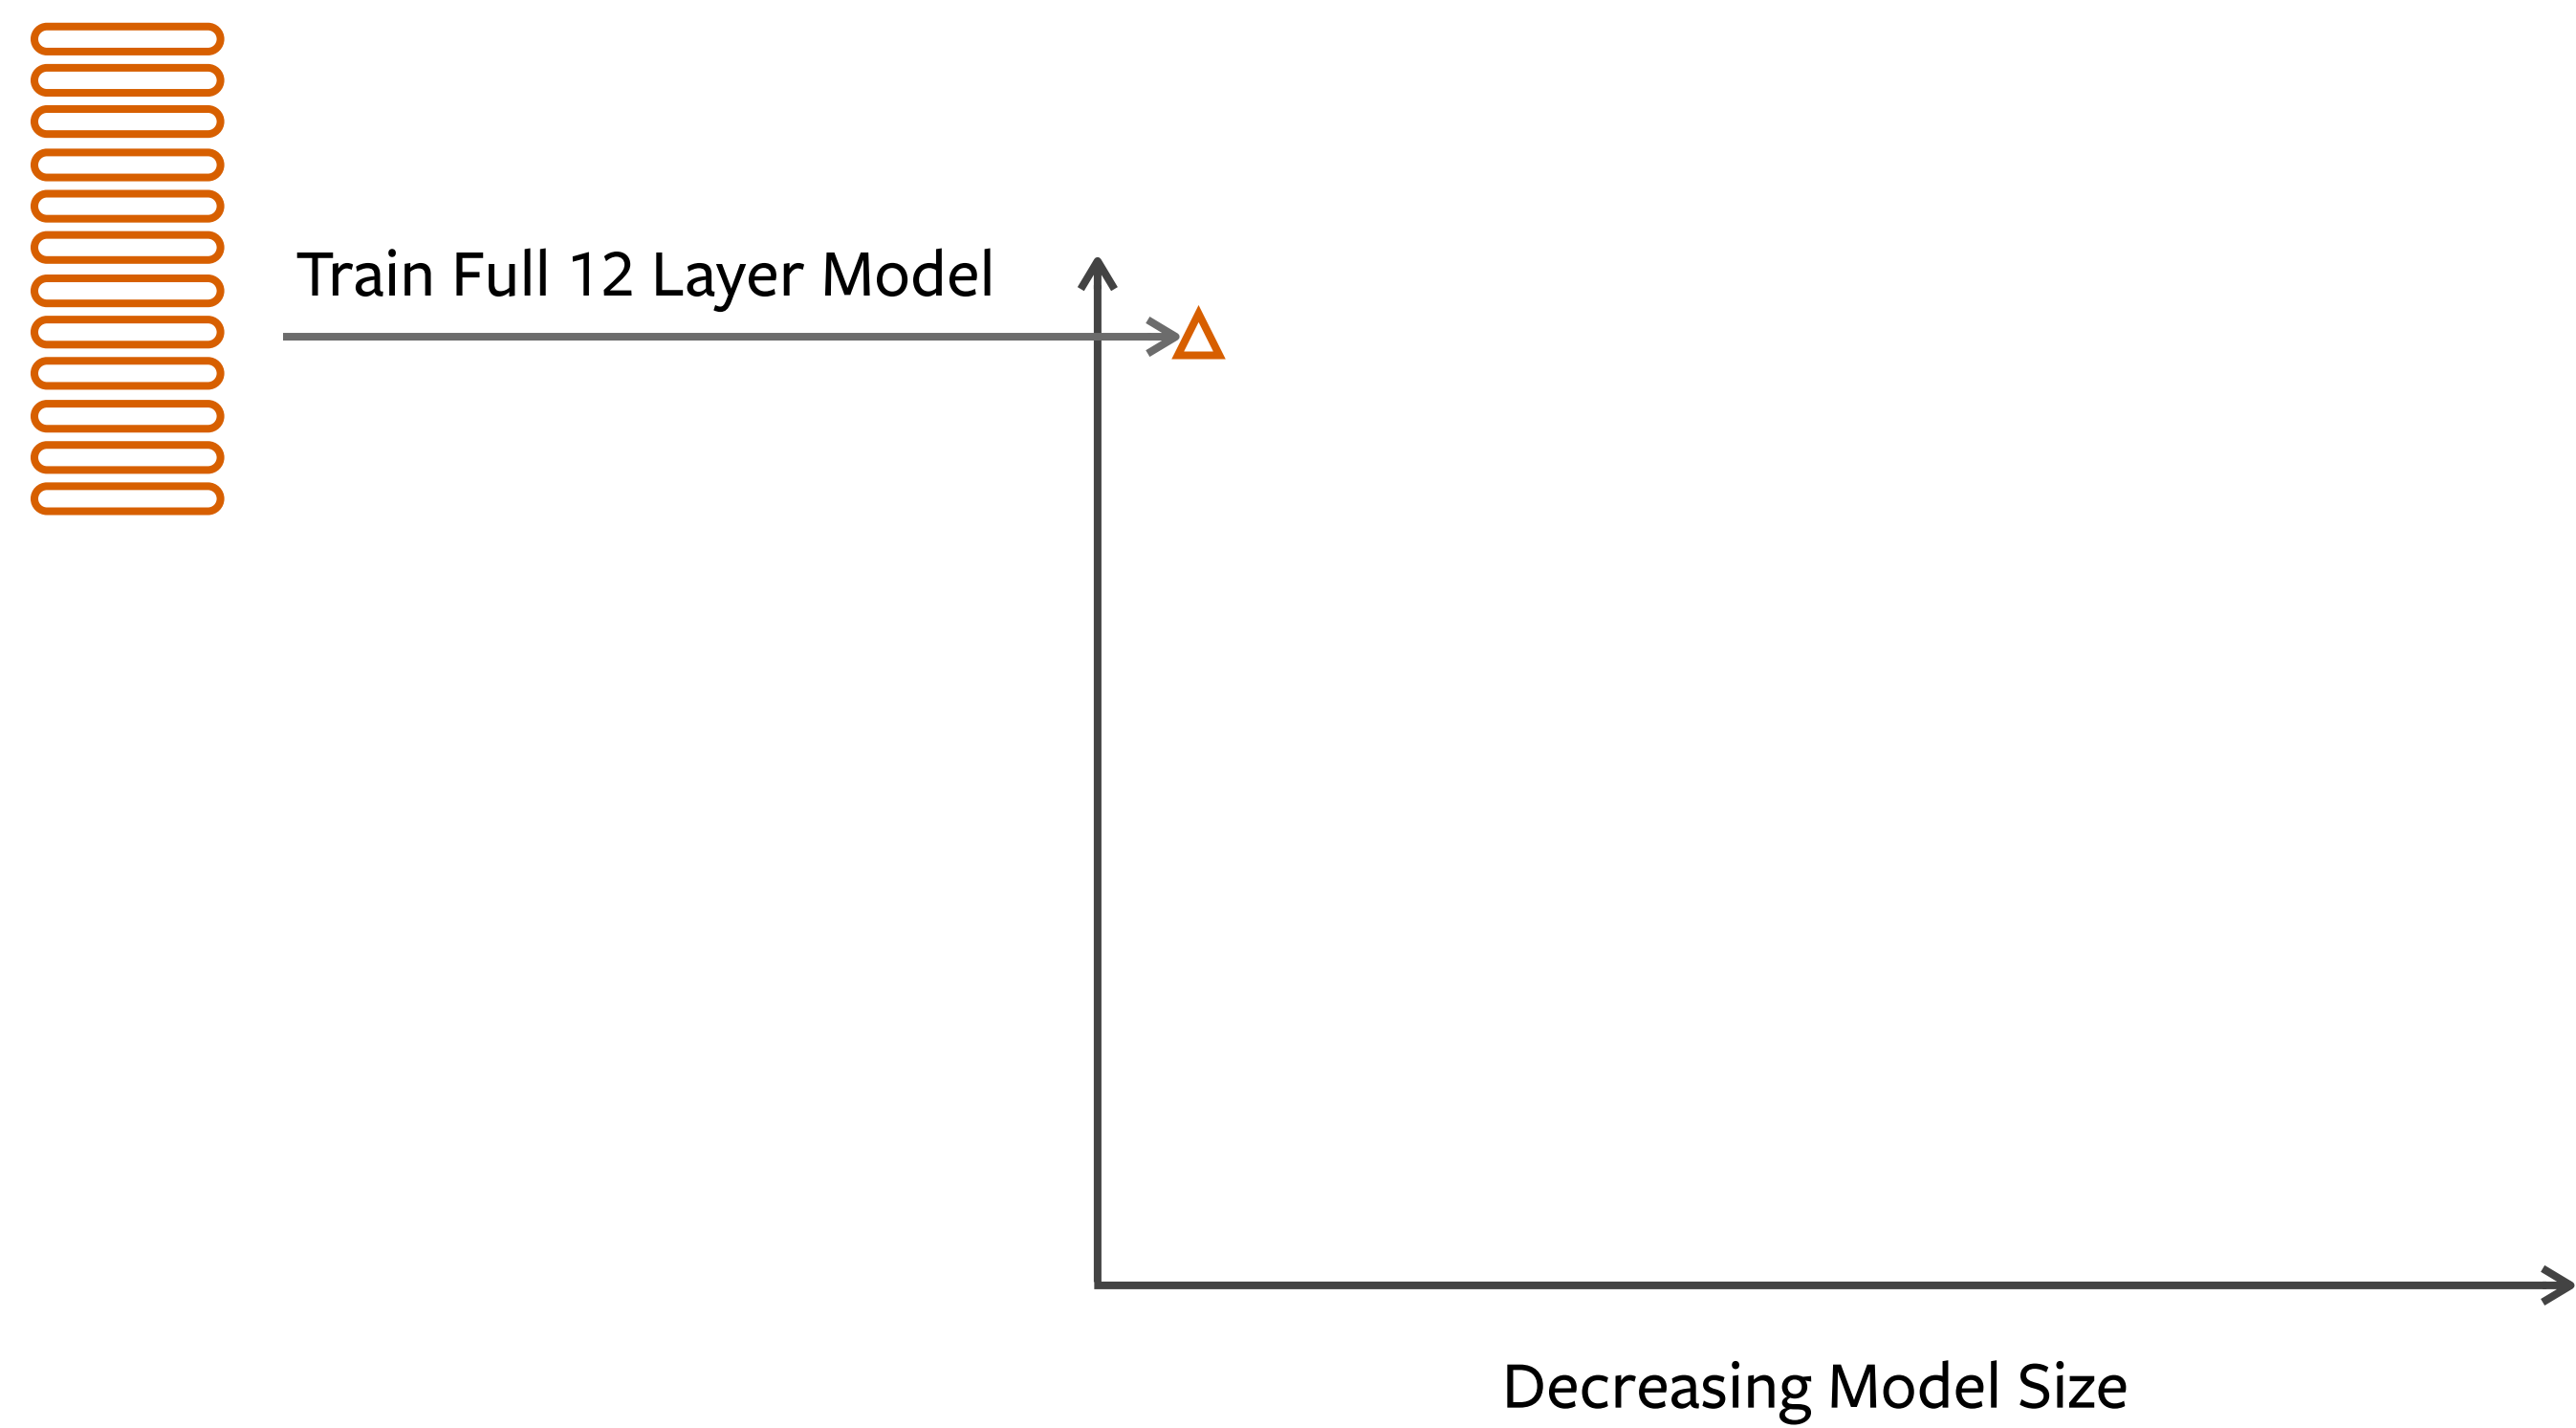
- Training Time
- Inference Time
- Model size
- Energy

not all techniques improve all of these areas


Techniques for Smaller Networks

- **Train Smaller Network from Scratch**
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization
- More efficient architectures

Training a Smaller Model from Scratch



Training a Smaller Model from Scratch



Training a Smaller Model from Scratch


- Training Time 
- Inference Time 
- Model size 
- Performance 

Techniques for Smaller Networks


- Train Smaller Network from Scratch
- **Sparsity Inducing Training**
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization
- More efficient architectures

Sparsity Inducing Training

NEURAL NETWORK WEIGHT MATRIX



Sparsity Inducing Training




sparsify the weight matrix to include many zeroes

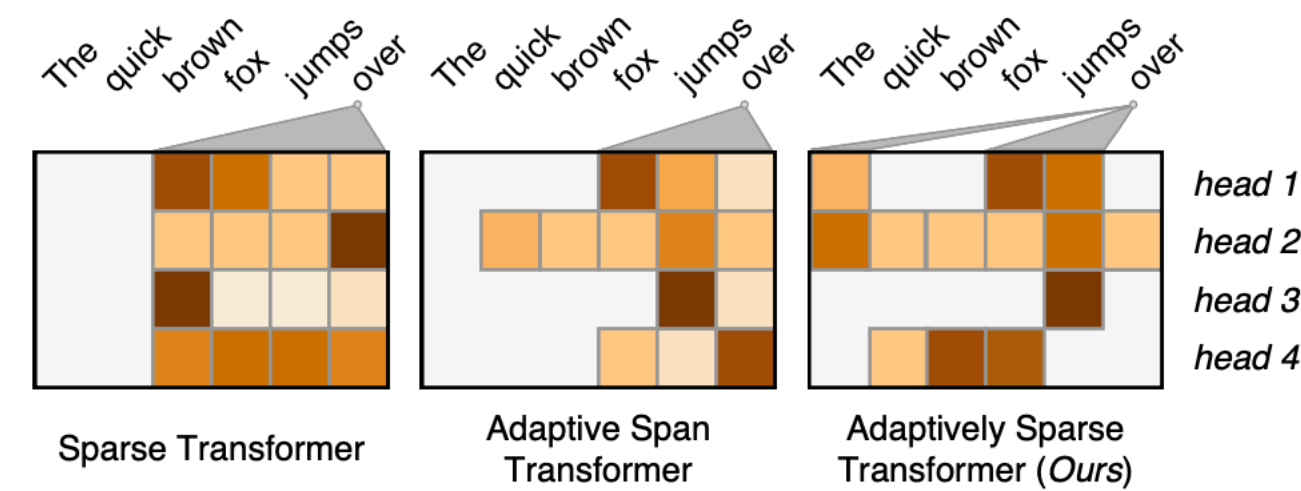
Advantages

- Sparse matrix multiplication - takes advantage of the zeroes
- Specialized kernels - everywhere you see a zero, don't need to compute that row/column, so less multiplications
- Important for on-device


Sparsity Inducing Training via Attention Matrices



GENERATING LONG SEQUENCES WITH SPARSE TRANSFORMERS
CHILD ET AL




ADAPTIVELY SPARSE TRANSFORMERS
CORREIA ET AL



STRATEGIES FOR STRUCTURING STORY GENERATION
FAN ET AL

Sparsity Inducing Training via Network Losses

penalize network for using parameters by increasing loss



ADAPTIVE ATTENTION SPAN IN TRANSFORMERS
SUKHBAATAR ET AL

Sparsity Inducing Training

- Inference Time



- Energy



- Performance



at least, often no performance drop



Techniques for Smaller Networks

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- **Knowledge Distillation**
- Pruning
- Weight Sharing
- Quantization
- More efficient architectures

Knowledge Distillation

LARGE, ACCURATE NEURAL NETWORK

SMALLER NEURAL NETWORK




TEACHER MODEL

STUDENT MODEL

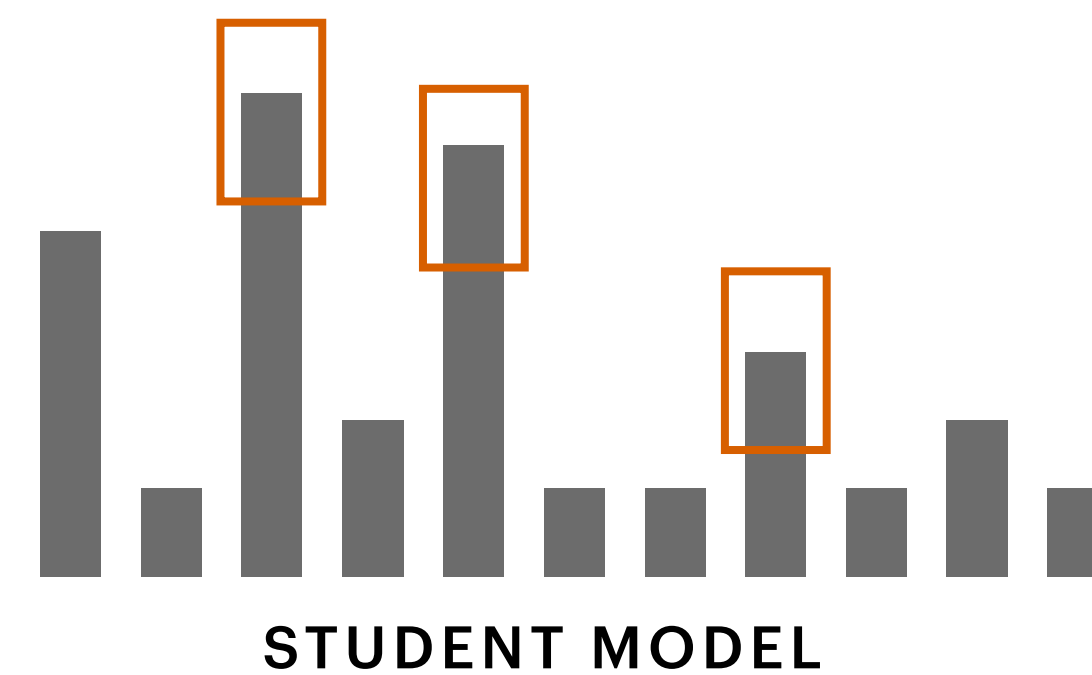

Knowledge Distillation

student model learns to mimic the output of the teacher model



Knowledge Distillation

student model learns to mimic the output of the teacher model



CALCULATE DIFFERENCE BETWEEN THE TWO PREDICTIONS

Advantages


- Flexibility over size - teacher and student can both be any size

Advantages

- Flexibility over size - teacher and student can both be any size
- Not limited to the training data - any data can be used to distill
 - data augmentation is **very powerful**

Advantages

- Flexibility over size - teacher and student can both be any size
- Not limited to the training data - any data can be used to distill
- Can also learn from intermediate layers



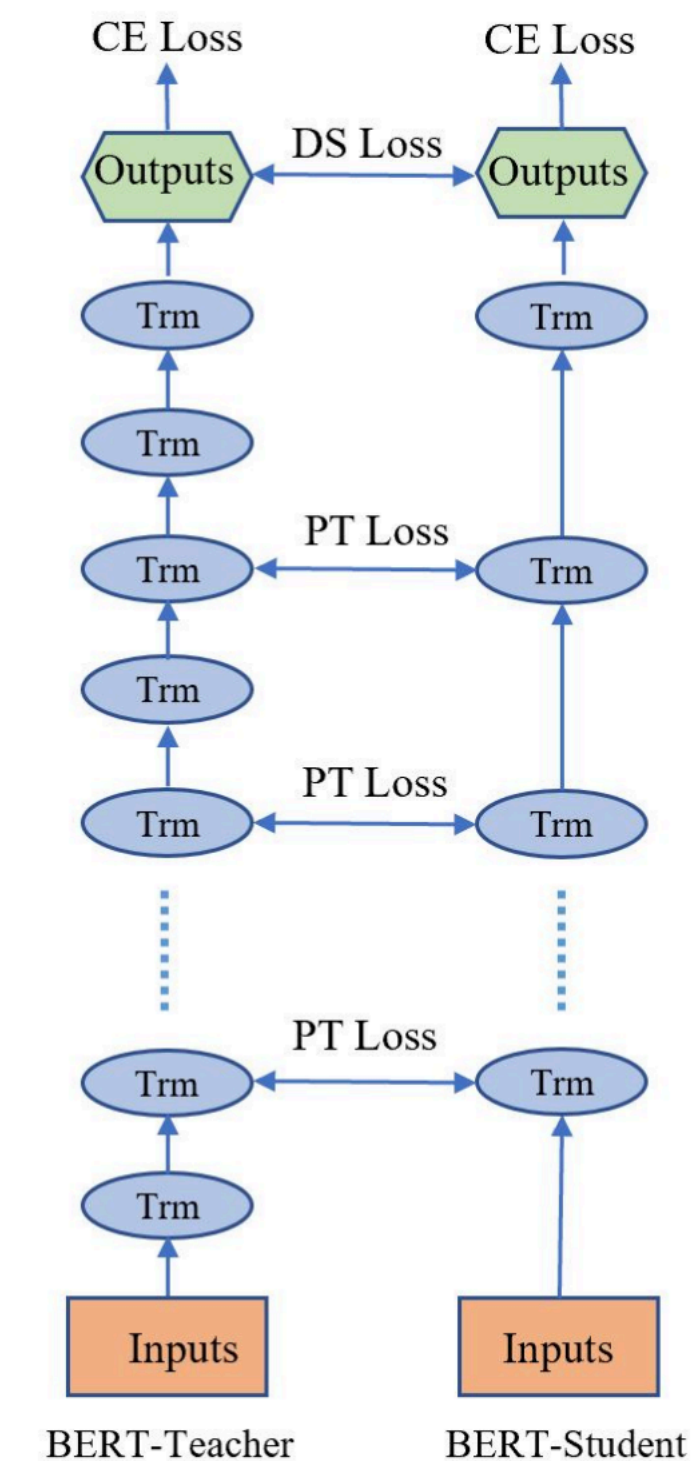
PATIENT KNOWLEDGE DISTILLATION FOR BERT MODEL COMPRESSION

SUN ET AL

Advantages

- Flexibility over size - teacher and student can both be any size
- Not limited to the training data - any data can be used to distill
- Can also learn from intermediate layers

PATIENT KNOWLEDGE DISTILLATION FOR BERT MODEL COMPRESSION
SUN ET AL



Knowledge Distillation

- Training Time
- Inference Time
- Performance



training time: need pre-trained teacher, but often
state of the art large models can be downloaded

Knowledge Distillation

people love knowledge distillation!

TINYBERT: DISTILLING BERT FOR NATURAL LANGUAGE UNDERSTANDING

JIAO ET AL

DISTILBERT, A DISTILLED VERSION OF BERT: SMALLER, FASTER, CHEAPER AND LIGHTER

SANH ET AL

WELL-READ STUDENTS LEARN BETTER: ON THE IMPORTANCE OF PRE-TRAINING COMPACT MODELS

TURC ET AL

MOBILEBERT: TASK-AGNOSTIC COMPRESSION OF BERT BY PROGRESSIVE KNOWLEDGE TRANSFER

SUN ET AL

AND MORE!

Techniques for Smaller Networks

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- **Pruning** remove layers from a trained model
- Weight Sharing
- Quantization
- More efficient architectures

Techniques for Smaller Networks

- Train Smaller Network from Scratch
- Sparsity Inducing Training if you want a smaller model size, need to re-train
- Knowledge Distillation
- **Pruning (with LayerDrop)**
- Weight Sharing
- Quantization
- More efficient architectures

How can we make models smaller without re-training?

Goal: Train One Network, Prune to Any Depth at Inference Time


How can we make models smaller without re-training?

Goal: Train One Network, **Prune to Any Depth at Inference Time**



Drop Any Layer and Model Remains the Same

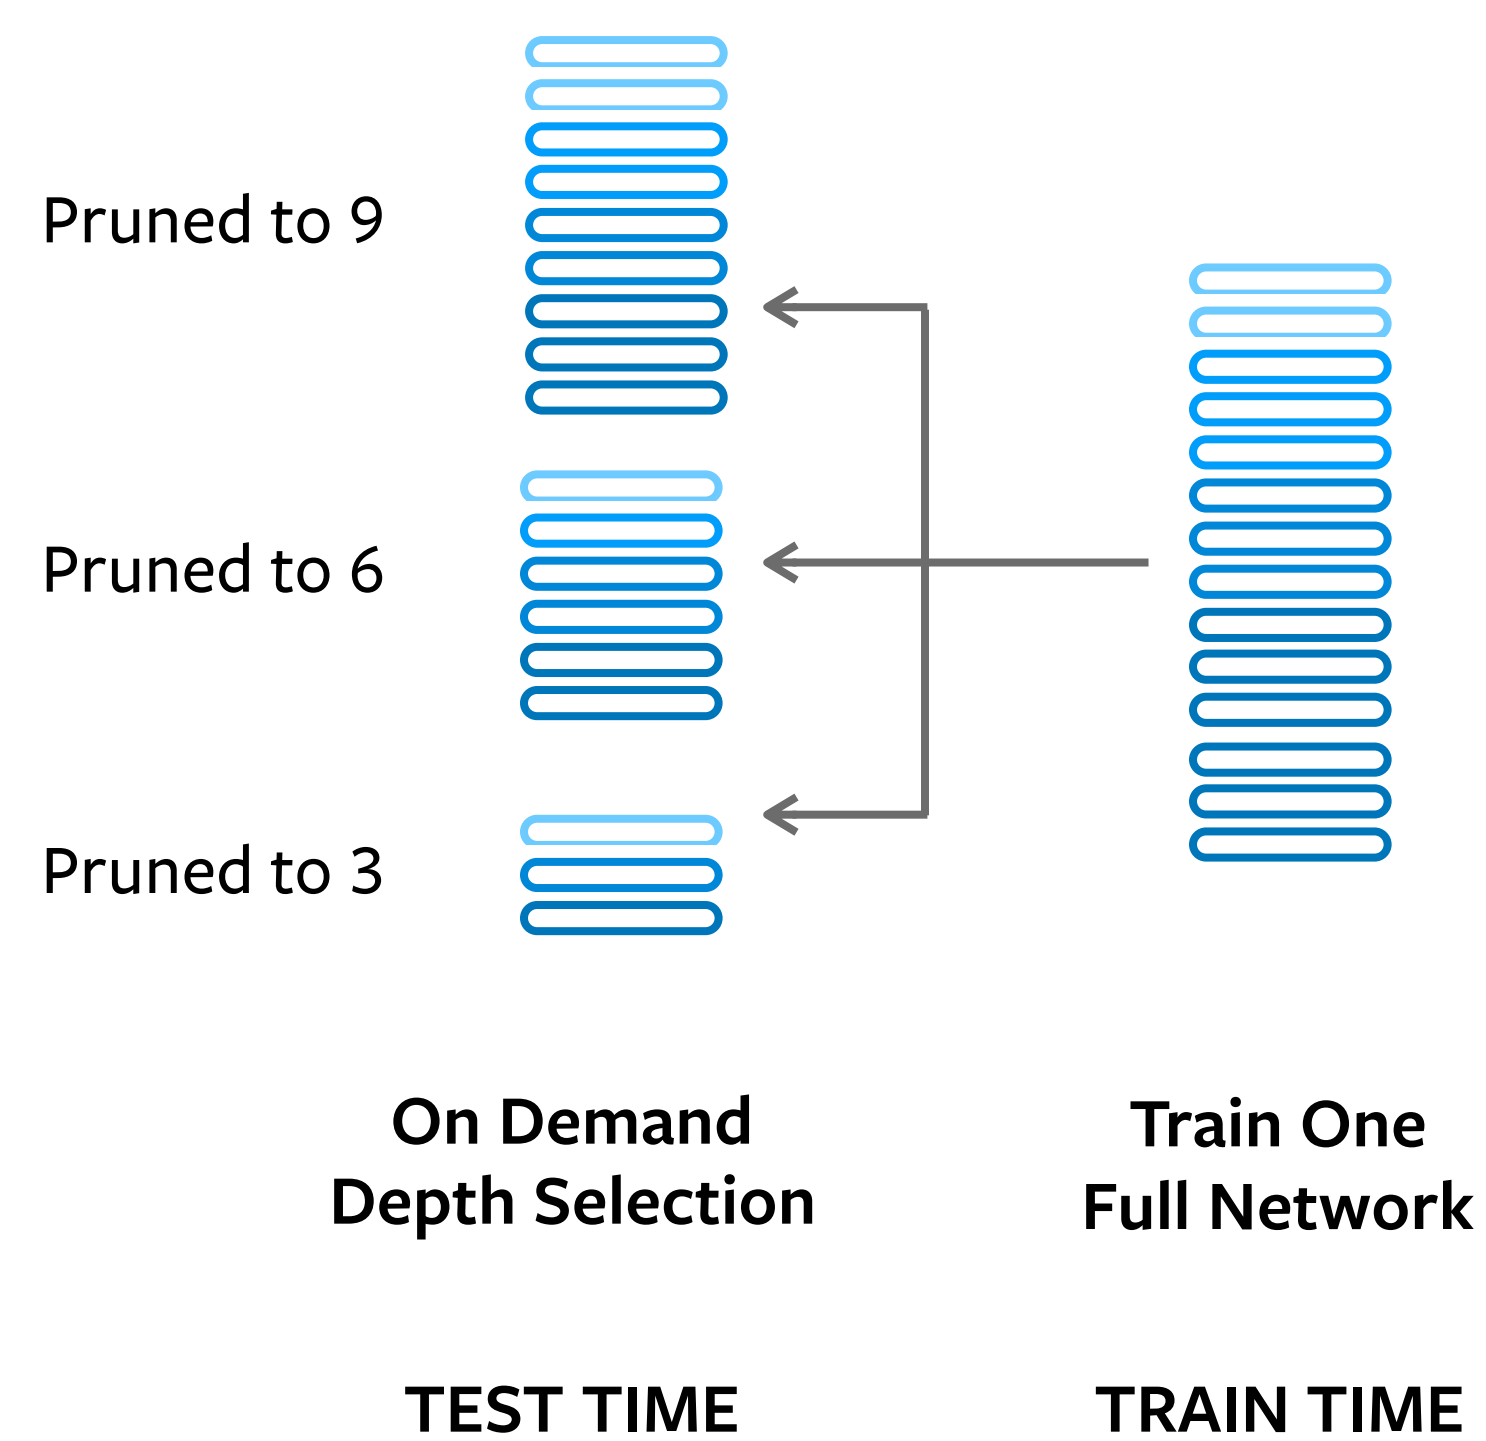
How can we make models smaller without re-training?



Train One
Full Network


TRAIN TIME

How can we make models smaller without re-training?




Our Proposal: LayerDrop

TRAINING TIME




Our Proposal: LayerDrop

TRAINING TIME




TEST TIME




Our Proposal: LayerDrop

TRAINING TIME




NO FINETUNING

TEST TIME




Structured Dropout can be More General

TRAINING TIME




TEST TIME




- Layers (our focus)

Structured Dropout can be More General

TRAINING TIME



TEST TIME




- Layers
- Attention Heads
- FFN, Attention Sub-Layers
- Portions of Matrices
- etc


Advantages

- Training Speed
- Regularization
- Reduction


(1) LayerDrop Increases Training Speed




(2) LayerDrop is an effective regularizer - Neural Machine Translation



(2) LayerDrop is an effective regularizer - Neural Machine Translation



(2) LayerDrop is an effective regularizer - Neural Machine Translation



(3) Our Main Focus: LayerDrop for Pruning

- Train once, prune to any desired depth


(3) Our Main Focus: LayerDrop for Pruning

- Train once, prune to any desired depth
- Robust to parameter setting
 - use the same value for all experiments


(3) Our Main Focus: LayerDrop for Pruning

- Train once, prune to any desired depth
- Robust to parameter setting
- Specific Pruning Strategy is not Important


(3) LayerDrop for Pruning - Language Modeling



(3) LayerDrop for Pruning - Language Modeling




(3) LayerDrop for Pruning - BERT Pre-Training




Train BERT
+ LayerDrop


(3) LayerDrop for Pruning - BERT Pre-Training




(3) LayerDrop for Pruning - BERT Pre-Training




(3) LayerDrop for Pruning - BERT Pre-Training




(3) LayerDrop for Pruning - BERT Pre-Training



(3) LayerDrop for Pruning - BERT Pre-Training




(3) LayerDrop for Pruning - BERT Pre-Training




Different Pruning Strategies - Does it Matter?

TRAINING TIME



TEST TIME



Add LayerDrop to Your Transformer Training

```
for layer in transformer.layers:  
    x = layer(x)
```

Add LayerDrop to Your Transformer Training

```
for layer in transformer.layers:  
    if random(0,1) > layer_drop and self.training:  
        x = layer(x)
```

Pruning with LayerDrop

- Training Time
- Inference Time
- Performance
- Model Size




Techniques for Smaller Networks

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- **Weight Sharing**
- Quantization
- More efficient architectures

Weight Sharing


NEURAL NETWORK



re-use weights in
multiple places

Weight Sharing

NEURAL NETWORK




input and output
embeddings tied
(common)

EXTREME LANGUAGE MODEL COMPRESSION WITH
OPTIMAL SUBWORDS AND SHARED PROJECTIONS
ZHAO ET AL

Weight Sharing

NEURAL NETWORK




share the weights of
chunks of layers

FAN ET AL

Weight Sharing

NEURAL NETWORK




share the weights of
chunks ALL layers

ALBERT: A LITE BERT FOR SELF-SUPERVISED LEARNING OF
LANGUAGE REPRESENTATIONS
LAN ET AL

UNIVERSAL TRANSFORMER
DEGHANI ET AL

Weight Sharing

- Model Size
- Performance



unless you increase model size

Techniques for Smaller Networks

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- **Quantization**
- More efficient architectures

Advantages of Quantization

- Easily combined with existing techniques
 - you can quantize a pruned model, quantize a distilled model, etc

Advantages of Quantization

- Easily combined with existing techniques
 - you can quantize a pruned model, quantize a distilled model, etc
- Ships with PyTorch and Tensorflow
 - easy to apply


Advantages of Quantization

- Easily combined with existing techniques
 - you can quantize a pruned model, quantize a distilled model, etc
- Ships with PyTorch and Tensorflow
 - easy to apply
- Can offer drastic compression


How much Model Size do you want to decrease?

- Train Smaller Network from Scratch maybe model will be 2-4x smaller
- Sparsity Inducing Training probably less...
- Knowledge Distillation 2-10x smaller
- Pruning 2x smaller
- Weight Sharing 2-8x smaller
- **Quantization** **4-25x smaller**
50 - 100x in combination

How does Quantization offer such extreme compression?



How does Quantization offer such extreme compression?



Different Types of Quantization

- Scalar Quantization (**int8**, int4, binary)
- Vector Quantization (Product Quantization)

Scalar Quantization

4x compression from int8

8x compression from int4

32x compression from binary... so far not working for Transformers

Scalar Quantization

- Neural Networks are often stored in fp32. We save space by going to int8 or int4.
- Take real numbers and instead store them as integers with scaling factors


0.0269
real number

Scalar Quantization


- Neural Networks are often stored in fp32. We save space by going to int8 or int4.
- Take real numbers and instead store them as integers with scaling factors

$$\begin{array}{ccccc} 0.0269 & & 110 & & 2^{-12} \\ \text{real number} & = & \text{integer} & * & \text{scaling factor} \end{array}$$

How to apply to an entire matrix?



Calculate Scaling Factor across all values



Vector Quantization: Product Quantization

25x compression or more



Vector Quantization: Product Quantization

25x compression or more



Subvectors

Codewords

Vector Quantization: Product Quantization



can combine scalar and vector quantization

Quantization

- Inference Time



if scalar quantization

- Performance



depending on how compressed

- Model Size

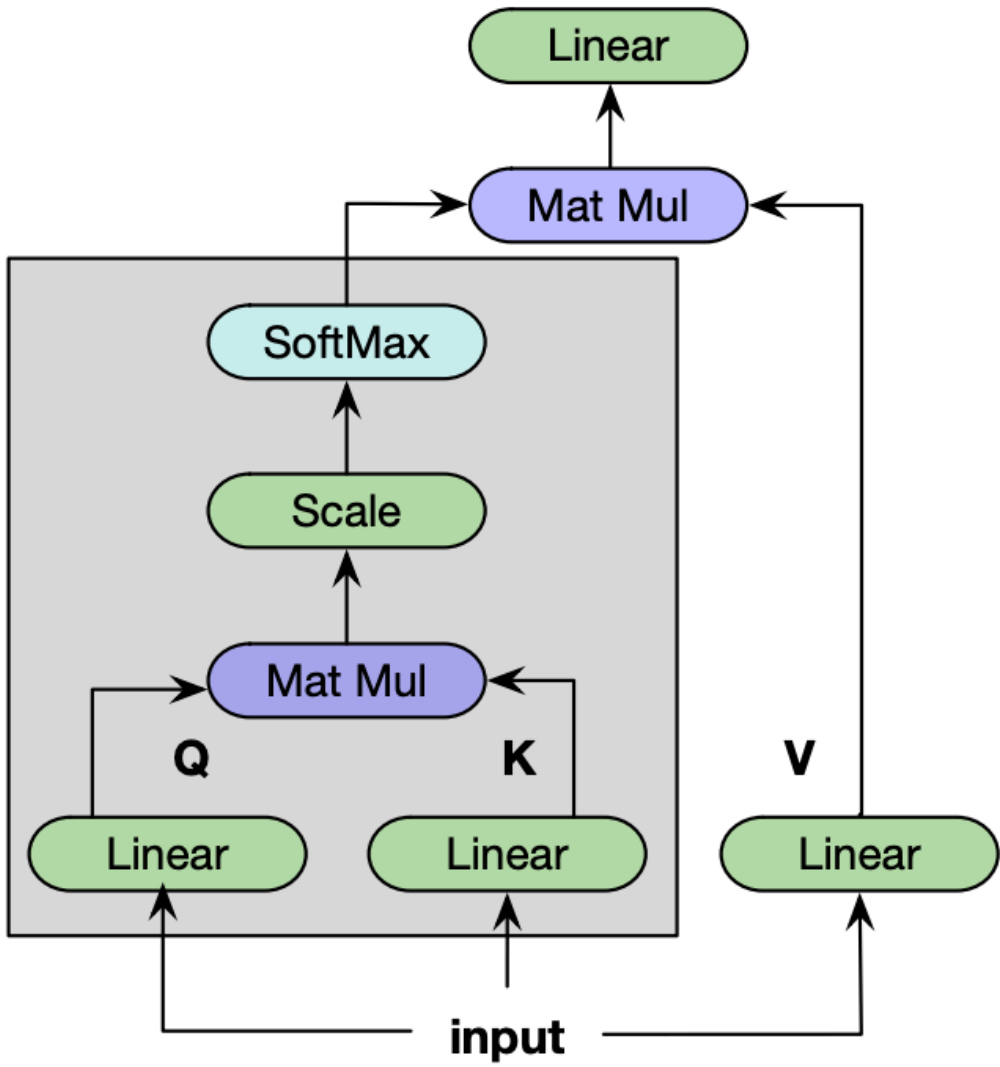


critical for on-device


Techniques for Smaller Networks

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization
- **More efficient architectures**

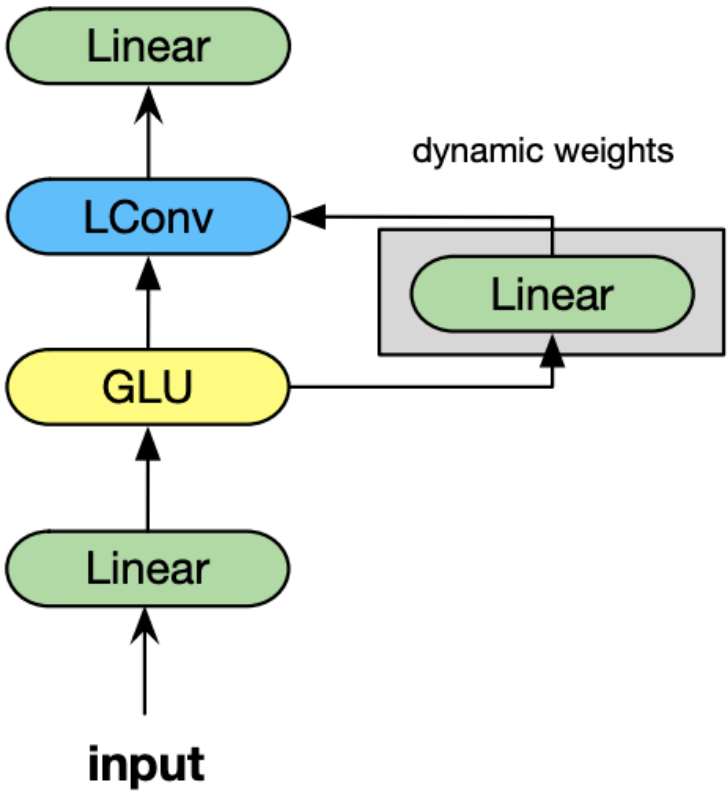
Variant Transformer Architectures



(a) Self-attention




(b) Lightweight convolution



(c) Dynamic convolution

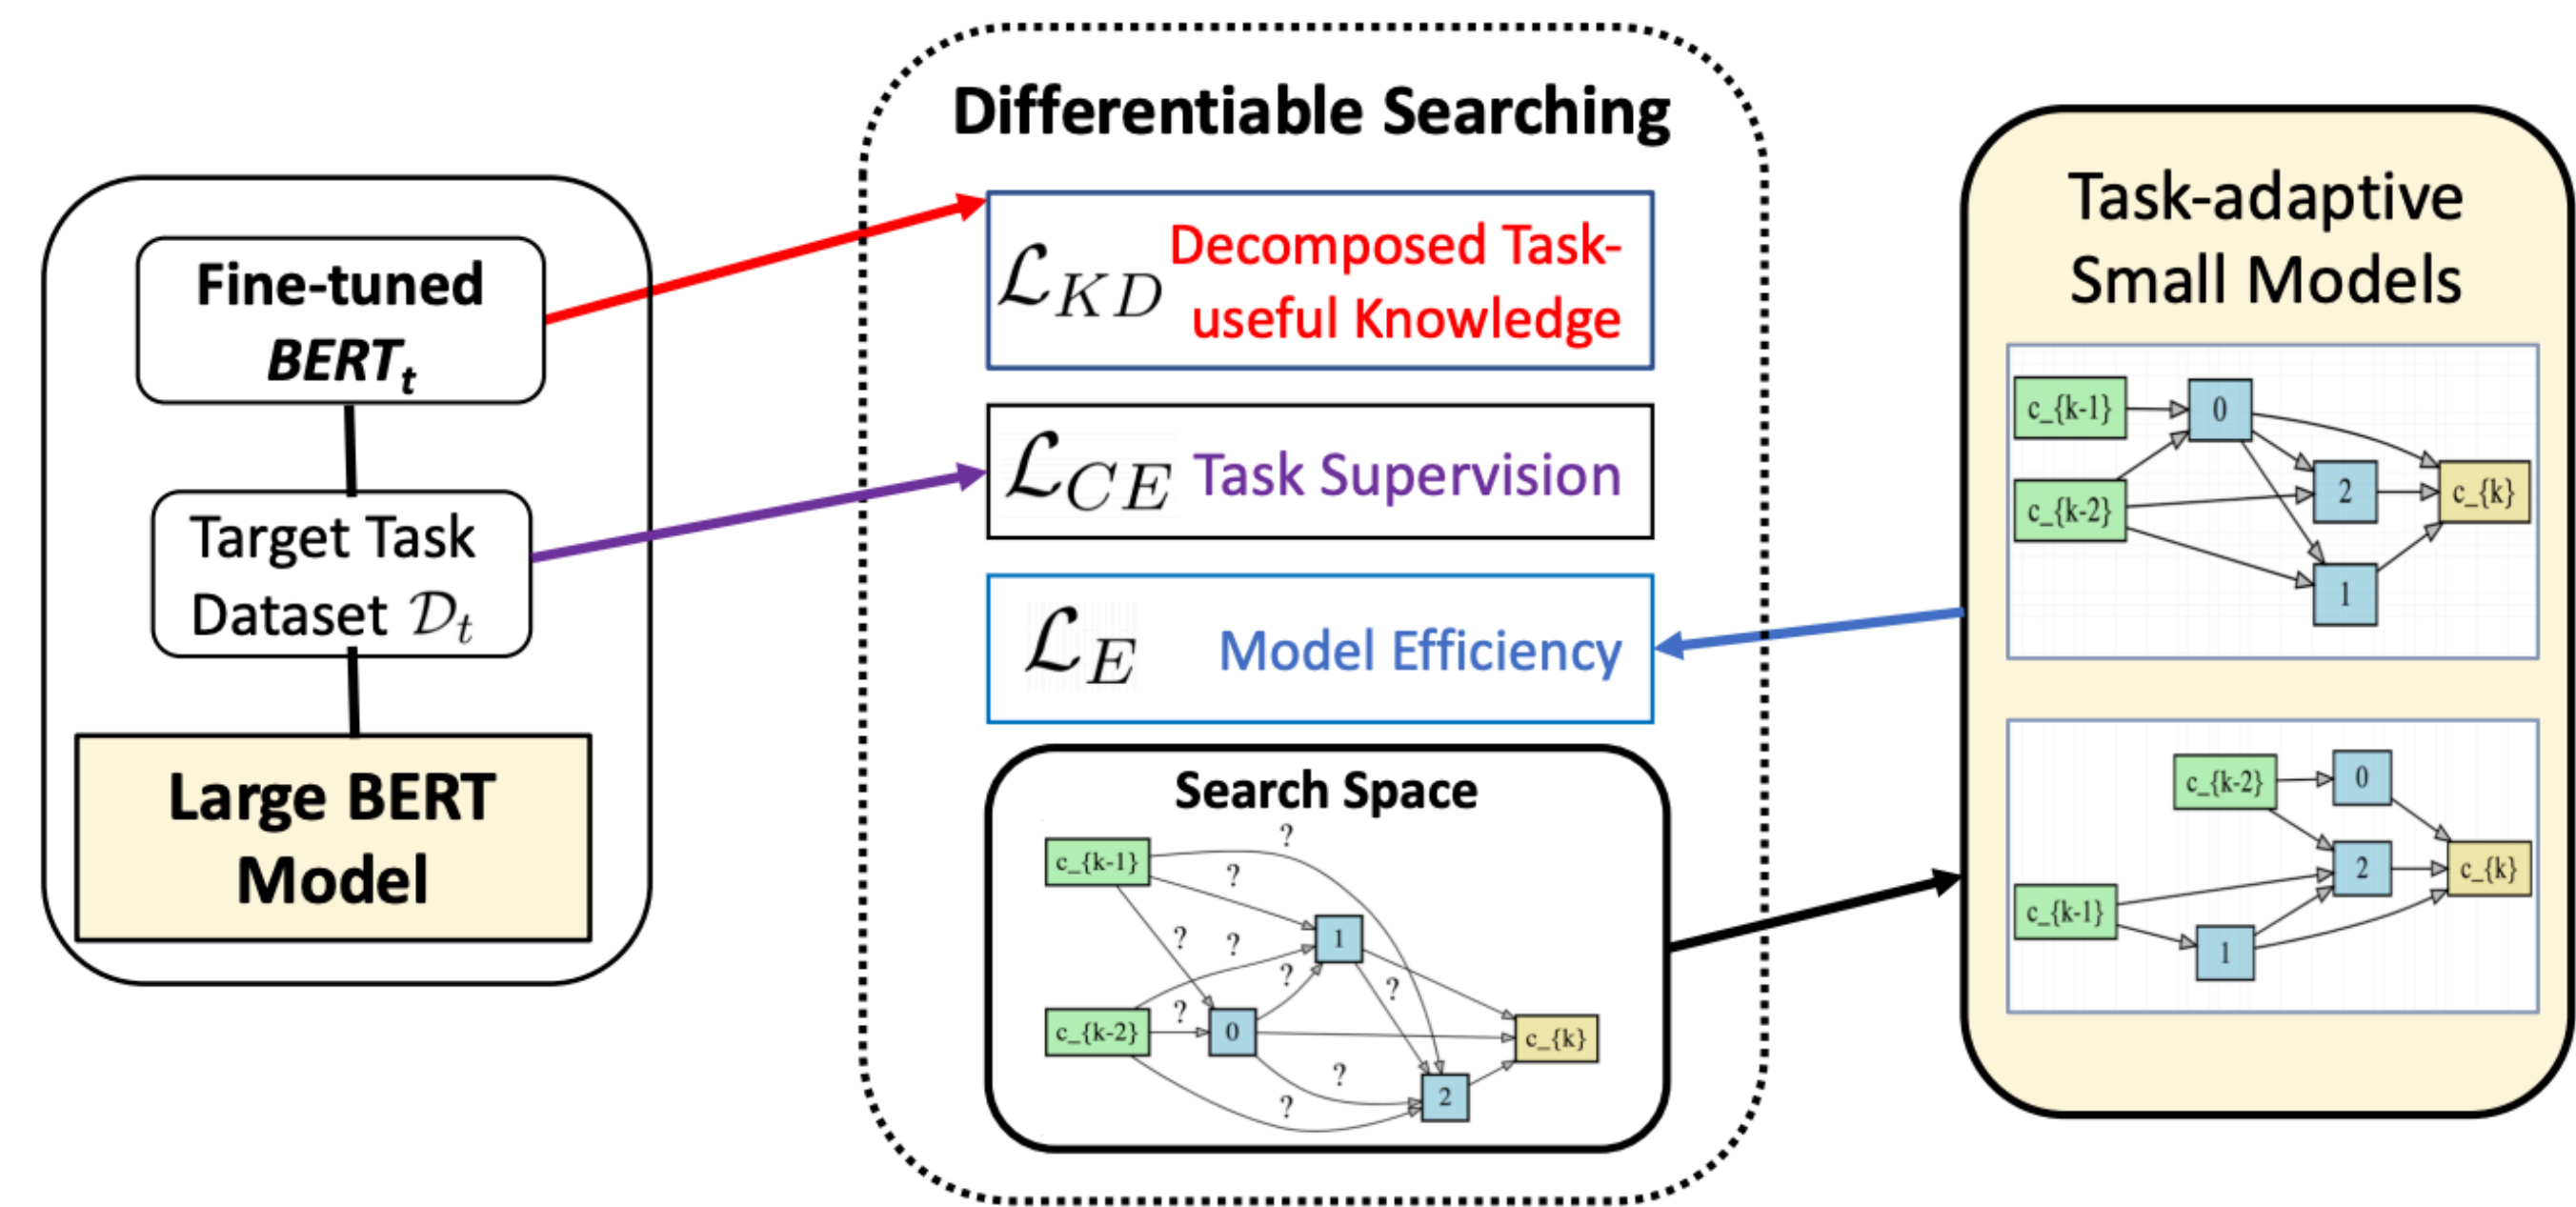
PAY LESS ATTENTION WITH LIGHTWEIGHT AND DYNAMIC CONVOLUTIONS
WU ET AL

Variant Transformer Architectures



ACCESSING HIGHER-LEVEL REPRESENTATIONS IN
SEQUENTIAL TRANSFORMERS WITH FEEDBACK MEMORY
FAN ET AL

Variant Transformer Architectures



ADABERT: TASK-ADAPTIVE BERT COMPRESSION WITH
DIFFERENTIABLE NEURAL ARCHITECTURE SEARCH
CHEN ET AL

In Summary...

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization
- More efficient architectures

strong baseline. would not discount.

In Summary...

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization
- More efficient architectures

important for on-device

In Summary...

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization
- More efficient architectures

a lot of recent improvements, very flexible

In Summary...

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization
- More efficient architectures

easy and straightforward gains

In Summary...

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization
- More efficient architectures

depends on how much you share

In Summary...

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization
- More efficient architectures

important for on-device. easily combinable.
can be used for aggressive compression

In Summary...

- Train Smaller Network from Scratch
- Sparsity Inducing Training
- Knowledge Distillation
- Pruning
- Weight Sharing
- Quantization
- More efficient architectures

exciting direction

And of course, even more considerations...

- Latency (faster decoding)
- Models that fit on Specialized Hardware
 - specific block sizes
 - battery life and heat from device

Interested in Efficient NLP?



Simple and Efficient Natural Language Processing

Workshop at EMNLP 2020 in Punta Cana

Thanks for listening!